

The Model Context Protocol Server Ecosystem: An Empirical Study of Quality, Security, and Evolution

Junan Chen^{a,*}

^aUniversity of Electronic Science and Technology of China, Chengdu, Sichuan, China

ARTICLE INFO

Keywords:

Model Context Protocol
software ecosystem
empirical software engineering
agent tooling
API documentation
credential management

ABSTRACT

AI agents reach external tools through the Model Context Protocol (MCP), and the official registry publishes every server version with registry-side timestamps, making a population-level study possible for the first time. We analysed a complete snapshot frozen on 24 September 2026: 114,604 published versions across 35,388 servers and 20,540 namespaces, enriched with package metadata for 13,316 packages, tool definitions extracted from 1,200 package sources, and repository data for 924 repositories. We established extraction quality against human coding of 150 packages (recall 99.7%, precision 99.7%), and two coders classified interface and security properties over 150 servers and 200 tool definitions. Four findings stand out. Publishing grew from 787 to 27,863 versions per month within twelve months, yet activity is concentrated: 100 of 20,540 namespaces account for 30.1% of all versions, and 62.7% of servers publish once and never update. Of 174,265 declared environment variables, 19.9% look like credentials but carry no secret flag, so a client cannot redact them automatically. Tool descriptions state purpose 93.5% of the time but selection guidance only 23.0% of the time, and 93.5% of coded tools are ambiguous with a sibling in the same server. Finally, 59.3% of servers that link a repository never tag a release, so registry version counts and repository release counts measure different things. The agent-tool layer is therefore only partly auditable from machine-readable artefacts, and we set out which properties are worth making machine-checkable.

1. Introduction

Large language model (LLM) agents act on the world through tools. In the last two years the Model Context Protocol (MCP) has become the usual way to expose those tools: a server declares a set of named operations with JSON-Schema parameters, a client discovers them at runtime, and the model chooses among them (Model Context Protocol, a). An official registry publishes every version of every server, so for the first time an entire agent-tool ecosystem is observable at population scale, with publication timestamps attached (Model Context Protocol, b).

That observability is also a liability, and one that software engineering has seen before. A registry is not a description of software quality; it is a description of what someone was willing to publish. The npm and PyPI ecosystems taught the community that registry metadata can be measured rigorously (Decan, Mens and Grosjean, 2019; Abdalkareem, Oda, Mujahid and Shihab, 2020), and that the measurements are often less flattering than the volume of packages suggests. The agent-tool layer deserves the same treatment before it hardens into infrastructure that other systems depend on.

This paper asks what an engineer can and cannot determine about an agent tool before trusting it, using only machine-readable artefacts. We take a complete snapshot of the official MCP registry, enrich it from package registries and repositories, extract tool definitions from published sources, and validate every extraction and coding step against human judgement. Our aim is not a laudatory

census. It is to establish which properties of the ecosystem are actually observable, and which are invisible.

1.1. Research questions

RQ1 Interface design. What interface design patterns do published MCP servers exhibit, in tool granularity, naming, description quality and input declaration?

RQ2 Security posture. How do servers handle credentials and remote exposure, and where does the declared configuration diverge from safe practice?

RQ3 Evolution. How does the ecosystem evolve: publication cadence, version churn, abandonment and protocol adoption?

RQ4 Inspectability. How much of the ecosystem can be assessed from machine-readable artefacts alone, and what does the remainder hide?

RQ4 is the study's methodological centre. It converts the census into a claim about auditability: how much of an agent-tool ecosystem resists automated inspection, and what would have to change for governance to be mechanisable.

1.2. Contributions

1. A complete, reproducible snapshot of the MCP server ecosystem, covering 114,604 published versions across 35,388 servers, with a validated pipeline that regenerates every table in this paper from public sources (§3).
2. A measurement of extraction reliability against human ground truth: 99.7% recall and 99.7% precision on 150 packages, with the reasons for every miss (§4.4).

*Corresponding author.

Email address: chenjunan.aiden@gmail.com (J. Chen)

ORCID(s):

3. The first population-level description of MCP interface quality, credential handling and evolution, each with inter-rater agreement reported per field (§4).
4. Evidence that registry version counts and repository release counts measure different things (Spearman 0.379), which constrains what any future study of this ecosystem can claim about release cadence. We generalise this into a proposition about registries as such: a registry can *directly* enforce the *shape* of a declaration, because shape is checkable at submission, but enforcing the *practice* behind the declaration needs post-hoc mechanisms, which act later and cost more. The distance between the two is measurable (§5.2).
5. An explanation, grounded in the coding results, of why agent-facing documentation converges where it does. Coding fields that ask whether a feature is present reach moderate agreement or better; fields that ask whether something is absent, sufficient or implied collapse to fair or slight. The same asymmetry is consistent with the ecosystem documenting naming thoroughly and selection conditions hardly at all (§5.3).

Below, Section 2 positions the work. Section 3 describes the data, sampling and coding. Section 4 reports results for RQ1–RQ4. Section 5 discusses implications. Section 6 states the threats to validity, including those we tested empirically. Section 7 concludes.

2. Background and Related Work

2.1. Agent tooling and the Model Context Protocol

An LLM agent selects tools at runtime from descriptions supplied by the server. The description is therefore part of the program: it is the only signal the model has when choosing between operations that sound alike. Catalogue and pattern work on foundation-model agents has begun to formalise how such systems are assembled (Liu, Lo, Lu, Zhu, Zhao, Xu, Harrer and Whittle, 2025a), and the security community has examined what happens when those descriptions turn hostile, through indirect prompt injection in tool-using agents (Khan, AlKhanbashi and Mohamed, 2026) and internal-channel leakage in multi-agent pipelines (Yagoubi, Badu-Marfo and Mallah, 2026). Both lines take the tool set as given. Neither asks what the tool set actually contains, or whether a client could have known before installing it.

MCP standardises the declaration itself: a server publishes a manifest describing how it is installed or reached, and exposes tools whose names, descriptions and input schemas are retrieved at runtime. An official registry lists published servers and retains every version rather than only the newest. Application papers have begun to report MCP integrations in specific domains, including adaptive traffic control (Othmani, Ben Yahia and Lalaguna, 2026), building energy modelling (Ball, Long, Fleming and Goldwasser, 2026) and orchestrated retrieval workflows (Zhou, Chen, Wang, Dong and Zhou, 2026). Each reports one server built

for one purpose. None measures the population those servers belong to, which is what this paper does.

2.2. Software ecosystem measurement

Researchers have measured package ecosystems rigorously for over a decade. Decan et al. compared dependency-network evolution across seven ecosystems and showed that growth statistics hide very different structural dynamics (Decan et al., 2019). Abdalkareem et al. showed that a large share of npm and PyPI packages are trivial, so package counts overstate capability (Abdalkareem et al., 2020). Work since has mapped the failure modes that volume conceals: dependency downgrades (Cogo, Oliva and Hassan, 2021b), same-day releases that flood dependents (Cogo, Oliva, Bezeemer and Hassan, 2021a), unused dependencies that persist in continuous integration (Weeraddana, Alfadel and McIntosh, 2024), the security cost of relying on small maintainers (Samaana, Costa, Abdellatif and Shihab, 2025), withdrawal and yanking (Li, Cogo and Bezeemer, 2023), and the structure of downstream package groups (Qi and Cao, 2024). Dependency bloat in JavaScript packages has been measured directly (Liu, Tiwari, Bogdan and Baudry, 2025b), as has the adoption of software bills of materials (Nocera, Romano, Di Penta, Francese and Scanniello, 2025).

Our study follows this line and differs in object. These works measure a dependency layer: what a package needs. We measure an interface layer: what a server offers and what it will do. The distinction matters because a model consumes that interface rather than a compiler, so ambiguity has consequences a dependency graph cannot express. We also follow the practice these works established of testing registry metadata against an independent signal, which §5.2 applies to version counts.

Two precedents deserve distinguishing. Large-scale characterisation of Java streams combined a new dynamic analysis with a population study of language-feature use (Rosales, Basso, Rosà and Binder, 2023); we lean on established analyses rather than new ones and invest instead in validating them. A study correlating README content with project meta-characteristics asked a related question on a single host (Venigalla and Chimalakonda, 2025); we ask it across two package registries and a protocol registry, and we measure what cannot be determined as carefully as what can.

2.3. API usability and documentation quality

That API design determines adoption and error rates is long established (Myers and Stylos, 2016; Robillard, 2009). Documentation is the part of that design a reader meets first, and its inadequacy has been documented from the developer side: documentation frequently omits exactly the information needed to decide whether and how to use an operation (Fan, Yu, Wang, Yin and Wang, 2020), and developers fall back on examples and source when prose fails them (Meng, Steinhardt and Schubert, 2019). Interface drift compounds the problem, since APIs evolve faster than their documentation (Hora, Robbes, Valente, Anquetil, Etien and Ducasse, 2018).

Table 1

Data sources and what each contributes. No source requires authentication except GitHub, where a read-only token is optional but raises the rate limit from 60 to 5,000 requests per hour.

Source	Role	Volume in this study
Official MCP registry	population of published servers, all versions, timestamps	114,604 versions / 35,388 servers
npm registry	package metadata, licences, download counts	9,557 packages
PyPI JSON API	package metadata and source distributions	3,759 packages
Package source archives	tool definitions, for static extraction	1,200 packages sampled
GitHub API	repository health, tagged releases	924 repositories sampled
Human coding	ground truth for extraction and interface quality	150 packages, 150 servers, 200 tools

These studies establish that documentation quality matters and that researchers usually assess it by human judgement. That raises a question they do not settle: how reliably such judgement can be applied. We address it in §3.5, reporting agreement per field rather than as a single pooled figure, because the fields do not behave alike.

2.4. Credential and supply-chain exposure

Automated dependency analysis is now standard, and security work has moved from package presence to package behaviour, up to detecting deliberately malicious PyPI packages (Umar Zeshan, Ibiyo, Di Sipio, Nguyen and Di Ruscio, 2026). Abandonment risk has become a signal in its own right and is scored explicitly for embedded components (Li, Moreschini, Pecorelli and Taibi, 2022); our RQ3 measures abandonment for the same reason, since a client deciding whether to depend on a server needs to know whether anyone still maintains it.

Credential handling in tool ecosystems has no comparable literature. Our RQ2 is a first measurement, and a narrow one at that: not whether a server is malicious, nor whether it mishandles secrets at runtime, but whether the artefacts it publishes tell a client enough to handle its credentials safely.

3. Study Design

This study measures one protocol ecosystem at one point in time, in cross section. Everything below is reproducible from public sources with the pipeline described in §3.7.

3.1. Data sources

Table 1 lists the sources. The registry comes first: it is the only place where the full population of published servers, all of their versions and registry-side timestamps coexist. We use package registries and repositories to test what the registry's own claims mean, never to expand the population.

3.2. Freeze window

Harvesting of the registry completed on 24 September 2026, and everything is frozen at that date. All statements of the form “as of the freeze” refer to it. Harvesting is resumable and every response is cached, so the snapshot can be reconstructed and verified: the SHA-256 digest of the raw harvest is recorded in the replication manifest.

3.3. Sampling

Two steps cost more than the population allows, so both use a seeded random draw:

- **Tool extraction** draws 1,200 packages from the 13,316 latest-version npm and PyPI packages the registry references (seed 20260924). Downloading every package would be tens of gigabytes.
- **Repository enrichment** draws 924 repositories from the 22,498 distinct repositories declared by published servers (seed 20260927).

We store both samples with their seed, so the analysed set is fixed. Neither draw takes the alphabetically first entries: repository ownership is not uniform across the alphabet, and an ordered cut would inherit that bias.

Why 1,200 packages. The sample size is set by the precision needed for the proportions we report, not by the download budget. A proportion estimated on n servers has a 95% confidence half-width of about $1.96\sqrt{0.25/n}$ in the worst case (a proportion near 0.5). At $n = 1,200$ that is ± 2.8 percentage points, which is smaller than every difference this paper reports: the gap between declared-share and credential-like share of environment variables is 18.7 points, and the gap between purpose coverage and selection coverage among tool descriptions is over 70 points. Halving the sample to 600 would widen the interval to ± 4.0 points and gain nothing, while tripling it buys precision we do not need and cannot use, because the binding constraint on the tool-level results is not the package count but how many sampled packages actually expose tools: 56.2% in this draw, or 675 servers.

The repository sample answers a different question, so it has a different size. It exists to test whether registry version counts track an independent release signal, which is a comparison of two measurements rather than an estimate of a population share. 924 repositories from a population of 22,498 give a half-width of about ± 3.2 points on the headline share, which is adequate for a test whose observed effect is a 59.3-point gap.

3.4. Tool extraction

Tool definitions live in implementation, not in the registry manifest, so recovering them means reading package

source. Both official SDKs register tools through a small number of idioms, and we match those idioms directly: a call to `tool` or `registerTool` on a server object, a tool object literal with `name` and `description` keys, a Python decorator on a function, or a `Tool` constructor call. Each match yields the tool name, the description where one is present, and whether an input schema appears inline.

Two worked cases show where this works and where it does not.

A clean match. The npm package `@aioproductoscom/mcp` declares three tools through object literals. Scanning nine source files recovers all three with their descriptions intact, for example:

triage, described as “Triage: turn fresh customer signal into prioritized, spine-linked work.”

A clean failure. The PyPI package `luma-mcp` registers four tools through decorators, but the decorator is applied indirectly and the matching expression never fires. The pipeline scans nineteen files and reports zero tools. Human inspection, invoked during validation, finds `get_event`, `search_events`, `export_event_ics` and `set_preferences`. This case is one of the five tools the validation pass charged against the pipeline, and it is the reason recall is reported as an estimate with its failure cases named rather than as a property of the method.

Failures cluster by language. Every missed tool and every over-counted non-tool in the validation sample came from Python packages, where decorator syntax admits indirection that literal matching cannot follow. We tested the obvious remedy: replacing pattern matching with an abstract syntax tree parse for Python. Re-scoring against the same human ground truth showed recall *falling* from 99.7% to 95.3%, because the two readers have complementary blind spots rather than one dominating the other. The validated configuration is the one reported throughout.

3.5. Measurement

Three quantities resist reading off a manifest, and we coded each by hand against a written codebook and an evidence pack:

- **Extraction validation** (150 packages): is the package actually a server, and how many tools does it expose? We established ground truth by reading the source rather than by trusting the pipeline.
- **Security posture** (150 servers, stratified by transport): how a client authenticates, whether a credential is required, how broad the tool scope is, and whether the tools can write or destroy.
- **Interface quality** (200 tool definitions, stratified by language): whether the description states purpose, selection condition, inputs and side effects, and which sibling tools an agent could confuse it with.

The evidence pack. Coding agent-facing interfaces fails when the coder cannot see the artefact being judged. Our first inter-rater round illustrates the failure mode precisely: the sheet given to the second coder carried an item identifier and a target name, and nothing else. Every field they could not look up came back as “cannot determine”. That is a disagreement about the task design rather than about the interface.

The fix is a single self-contained document per coding round. For each item it reproduces the manifest fields the registry declares, the declared environment variables with their required and secret flags, the verbatim tool description, and the list of sibling tools in the same server. Where a rule is mechanical it also shows the rule’s verdict, so the coder confirms rather than derives. The effect is that fifty items across two parts became roughly an hour of work instead of an afternoon, and, more importantly, that the coding measures the artefacts rather than the coder’s ability to locate them.

Codebook revision as part of the method. The codebook went through three versions, each answered by a coding round, and the sequence is worth reporting because it is reproducible in a way the final rule set alone is not.

Version 1 produced near-zero agreement on every field. Diagnosing it separated two causes that look identical in the output: an incomplete handoff, described above, and two defective value sets. One offered both `unclear` and `n/a` for “the evidence does not settle this”, which one coder read as synonyms and the other as opposites; collapsing the two post hoc raised observed agreement on the affected field from 0.00 to 0.56, which confirmed the defect. The other had no rule for a field that cannot be answered from the collected artefacts at all, and the two coders defaulted in opposite directions.

Version 2 replaced both words with a single `undetermined`, gave `scope_breadth` the value its sample demanded, and introduced the evidence pack. Seven of eleven fields then reached moderate agreement or better; four did not, and each failure had a different cause.

Version 3 removed `n/a` from the side-effect field so that it measures disclosure rather than read-only status, restricted secret declaration to manifest evidence, anchored selection guidance with examples adjudicated from the actual disputes, and withdrew the one field neither coder could answer for any item.

Only the third round is reported as reliability. The earlier rounds and all three codebooks ship with the replication package, so the sequence can be reproduced rather than taken on trust.

A second coder independently coded 50 shared items in that final round. We report Cohen’s κ (Cohen, 1960) and observed agreement per field, interpreted with the bands of Landis and Koch (Landis and Koch, 1977), and we do not report a field as a point estimate where agreement is below 0.40. Table 8 gives the full table, and §5.3 reads the pattern of failure rather than treating it as a defect to be disclosed.

Table 2

Percentile bootstrap 95% intervals for the sampled quantities (10,000 resamples). Census figures are omitted by design: they are population values, not estimates.

Quantity	Point	95% interval
<i>Coded tool definitions (n = 200)</i>		
Purpose stated	93.5%	[90.0, 96.5]
Inputs explained	61.0%	[54.0, 67.5]
Side effects named	38.5%	[32.0, 45.0]
Selection condition stated	23.0%	[17.5, 29.0]
Confusable with a sibling	93.5%	[90.0, 96.5]
<i>Coded servers (n = 150)</i>		
Authentication identified	64.0%	[56.7, 71.3]
No authentication required	34.7%	[27.3, 42.7]
Write capability	34.0%	[26.7, 42.0]
Broad scope	31.3%	[24.0, 38.7]
Broad scope and write capability	22.0%	[15.3, 28.7]
<i>Sampled packages and repositories</i>		
Yielded at least one tool (n = 1,200)	56.3%	[53.6, 59.2]
No tagged release (n = 914)	59.3%	[56.1, 62.5]
Spearman agreement (n = 914)	0.379	[0.314, 0.442]
<i>Credential-like variables, clustered by server</i>		
Not flagged secret, of credential-like	44.1%	[40.3, 47.5]
Not flagged secret, of all declared	19.9%	[18.3, 21.6]

3.6. Statistical treatment

The distributions in this ecosystem are heavy-tailed, with a few publishers producing thousands of versions and most producing one. We therefore report medians and concentration alongside every mean, and we treat cross-source comparisons as descriptive associations rather than hypothesis tests. No *p*-values are reported.

We do report uncertainty, but only where it exists. The registry census covers the whole population of published servers and versions, so its shares carry no sampling error and are given as exact counts; attaching a confidence interval to a census would assert an error that is not there. The figures drawn from a sample do carry sampling error, and Table 2 gives percentile bootstrap intervals for them (10,000 resamples, seed 20260928). The credential shares resample *servers* rather than variables, because variables cluster within a server and resampling variables would understate the interval.

Two of those intervals are worth stating here because they bound claims made later. Selection guidance in tool descriptions is 23.0% with a 95% interval of [17.5, 29.0], so the ordering against purpose coverage at 93.5% [90.0, 96.5] is not in question even though the point estimate is. Rank agreement between registry versions and repository releases is 0.379 [0.314, 0.442], which excludes perfect agreement and also excludes independence.

The credential figures are quoted on all published versions; restricting to the latest version of each server gives 17.6% [16.7, 18.6] of declared variables, lower because servers that republish often repeat their declarations. We

report the all-versions basis throughout, since the object of study is the registry as published.

3.7. Pipeline

Seven harvest-and-analysis steps and three validation steps make up the pipeline; all are resumable and cached. Each step writes the tables reported here; no number in this paper is computed by hand. The replication package contains the code, the drawn samples, the generated tables and the manifest of SHA-256 digests.

4. Results

4.1. RQ1: Interface design

Of the 1,200 sampled packages, 1,169 (97.4%) parsed successfully and 675 (56.2%) yielded at least one tool definition, producing **12,065 tool definitions across 676 servers**, a mean of 17.8 tools per server. The remaining packages either are not servers in the sense of exposing tools (libraries, clients, meta-packages) or register tools in ways static analysis cannot see. §4.4 quantifies that split.

Naming is markedly conventional. Of the 35,388 published servers, 99.5% use a lower-case server segment, none omits the slash separator and none departs from a dotted namespace: the registry enforces the reverse-DNS shape. Within the extracted tools, 86.5% use *snake_case*, 3.2% *kebab-case* and 2.2% *camelCase*, so a model moving between servers meets a dominant convention rather than a mix.

Description coverage is high but shallow. 96.4% of the 12,065 extracted tools carry a description, yet 11.3% of those are shorter than five words, and only 3.2% of tools expose an input schema our extraction could recover. We treat that figure as a lower bound, because schemas are frequently passed as variables or generated, and therefore as a limitation of static reading rather than a property of the servers.

Table 3 reports the coding of 200 tool definitions. Descriptions reliably convey *purpose*: 93.5% state what the tool does (95% CI [90.0, 96.5]), at complete agreement (Table 8). They convey *selection guidance* far less reliably: 61.0% explain their inputs [54.0, 67.5], and only 23.0% state when an agent should choose the tool [17.5, 29.0], a figure qualified in §4.4.1. The interval earns its place here: the two coders disagreed about this field, yet both readings sit far below the purpose figure, so the ordering is not in question.

The most striking result is ambiguity. **93.5% of coded tools were judged confusable with at least one sibling in the same server** (95% CI [90.0, 96.5]), with a mean of 2.31 confusable siblings per tool. The ecosystem has settled on how to name tools and largely on how to describe them, but not on how to make them distinguishable from one another.

Python tools are better documented than JavaScript ones on inputs (74.6% against 55.3%) and on selection guidance (30.5% against 19.9%), while purpose coverage is equal to within a percentage point.

4.2. RQ2: Security posture

Published manifests declare **174,265 environment variables**. Publishers mark 18.0% of them required and 26.5%

Table 3

Interface quality, coded by hand over 200 tool definitions sampled across the two dominant languages (141 JavaScript/TypeScript, 59 Python).

Property	Share of 200 tools
Description states the purpose	93.5%
Description explains the inputs	61.0%
Description names side effects	38.5%
Description states when to use the tool	23.0%
Confusable with a sibling tool	93.5%

secret. Applying a name-and-description heuristic, 45.2% look like credentials (Fig. 1). That gap is the finding: **34,704 variables, 19.9% of all declared variables, look like credentials but carry no secret flag**, so a client has no instruction to redact them from logs, prompts or transcripts. 3,273 of them are also marked required, meaning a user must supply a value and still receives no signal that it is sensitive. A further 29.1% of declared variables carry a default value, which in a credential context means a value published in a manifest. The share this subsection turns on is 19.9% of declared variables (95% CI [18.3, 21.6]); restricted to the latest version of each server it falls to 17.6% [16.7, 18.6], because servers that republish often repeat their declarations. We quote the all-versions basis throughout.

Across the population the manifests declare 48,706 remote endpoints on 14,903 distinct hosts. Against expectation we found **no plaintext http endpoints, no localhost endpoints, and a single endpoint on a bare IP address**. This is a negative result that we attribute to the registry's own validation rather than to publisher discipline (§6), and we report it as evidence about the gatekeeper rather than about the ecosystem.

Table 4 reports the coding of 150 servers stratified by transport (84 remote-only, 57 package-only, 7 both, 2 neither). Authentication could be determined for 64.0% of servers and remained undetermined for 36.0%, a result about artefact completeness as much as about security.

Two patterns matter. First, a substantial minority require no authentication at all (52 servers, 34.7%), consistent with read-only public data services. Second, capability and scope concentrate: 51 servers expose at least one writing tool and 10 a destructive one, 47 have broad scope, and **33 servers, 22.0%, combine broad scope with write capability**, the combination that turns a confused or compromised agent into a consequential event. Seven servers gate access through payment rails (x402 or prepaid credit) rather than identity, a category we did not anticipate when writing the codebook.

4.3. RQ3: Evolution

Growth was steep across the twelve months in the snapshot: from 787 published versions in September 2025 to **27,863 in September 2026**, with 9,421 servers appearing for the first time in that final month alone. Cumulative servers reached 35,388, and the curve shows no inflection (Fig. 2).

Table 4

Security posture, coded by hand over 150 servers stratified by transport.

Property	Servers (of 150)
Authentication mechanism identified	96 (64.0%)
Authentication mechanism undetermined	54 (36.0%)
No authentication required	52 (34.7%)
Credential-based authentication	37
Payment-gated (x402, prepaid credit)	7
Requires the user to supply a secret	24
of those, secret not flagged	3
At least one writing tool	51
At least one destructive tool	10
Broad scope	47
Broad scope and write capability	33 (22.0%)

Table 5

Concentration of publishing activity across namespaces.

Namespaces	Share of versions
Top 1	2.1%
Top 5	8.4%
Top 10	13.2%
Top 25	18.8%
Top 50	24.0%
Top 100	30.1%
Bottom 50% (10,270)	9.0%

Growth is not distributed evenly. Table 5 shows that the 100 most prolific namespaces, 0.49% of the 20,540 namespaces, produce **30.1% of all published versions**, while the bottom half produce 9.0%. A single namespace published 2,387 versions, 2.1% of the population by itself. "Versions published" is the most convenient proxy for activity in a registry, and here it is dominated by a small number of automated publishers. Every activity statistic in this paper therefore carries a median beside its mean, and we state the concentration wherever a total appears.

62.7% of servers published exactly one version and never updated again, a further 16.1% last published within 30 days of the freeze, and only 0.2% had gone more than a year between their last two publications. 374 servers (1.1%) are formally deprecated. Fig. 3 shows how concentrated the release distribution is.

Package-based servers behave differently from hosted ones. Servers shipping an installable package average 5.19 versions against 1.77 for hosted-only servers, and declaring a repository is associated with 3.59 against 2.09. Deprecation follows the same split: 1.8% for package-based servers against 0.5% for hosted-only. The pattern is consistent with packages being versioned artefacts and hosted endpoints being deployed in place.

Transports shifted during the window. Counting versions, the local stdio transport still dominates at 67.0%, but the hosted streamable-http transport appears in 42.1% of versions across 21,208 servers, more than stdio's 15,009;

Table 6

Manifest field completeness across all 35,388 published servers.

Field	Present	Missing
description	35,388	0
version	35,388	0
repository	27,146	8,242
title	21,767	13,621
remotes	21,526	13,862
websiteUrl	17,386	18,002
packages	15,166	20,222

hosted servers therefore publish fewer versions each. The legacy sse transport is residual at 2.4%.

Packaging concentrates on npm (54.5% of package declarations), then PyPI (24.2%), the MCP bundle format (10.9%) and OCI images (9.1%). Manifest schema adoption is near-complete: 92.3% of versions declare the 2025-12-11 schema.

Quality signals around those packages are weaker. Of 9,557 enriched npm packages, **30.1% declare no licence**, 67.4% have a single maintainer, 2.1% are deprecated, and 15.5% recorded no downloads in the last month (median 254). PyPI is worse on licensing: **85.0% of 3,759 PyPI packages declare no licence**, with 1.1% carrying yanked releases.

Repository health points the same way. Of 751 resolved repositories, 20.4% declare no licence, 51.7% have no stars and only 2.8% have 100 or more, 5.6% have not been pushed in 180 days, and the median repository is 114 days old. This ecosystem is young, distributed and almost entirely unstarred.

4.4. RQ4: Inspectability

Table 6 reports field completeness across all 35,388 servers. Description and version are always present because the registry requires them. Everything that would let an auditor act is optional: **23.3% of servers link no repository**, 39.2% declare no remote endpoint, 57.1% declare no package and 50.9% give no website. A quarter of the population cannot be traced to source code at all.

Our pipeline reads as much as an automated reader reasonably can: it downloads published package sources and parses the registration idioms of both official SDKs. It parsed 97.4% of the sampled packages and recovered tools from 56.2%. The complement, 43.8% of packages yielding no tool definition, is the inspectability gap. Human validation on 150 packages (Table 7) shows the gap decomposes into two unequal parts: 69 of the 150 are not tool-exposing servers at all, and only a small residue is genuine extraction failure.

Where the pipeline did find a server it found essentially all of it: **99.7% recall and 99.7% precision** on 1,794 hand-counted tools, with five tools missed across two packages and five non-tools over-counted across three. Every miss and every false positive falls in a Python package, where pattern-based extraction of decorator-registered tools is weakest. We take this seriously enough to have tried the obvious remedy:

Table 7

Human validation of tool extraction over 150 packages.

Metric	Value
Packages classified as a server	81
Servers where the extractor found at least one tool	78
Package-level detection rate	96.3%
Tools counted by hand	1,794
Tools matched by the extractor	1,789
Estimated tool-level recall	99.7%
Estimated tool-level precision	99.7%

Table 8

Inter-rater agreement over 50 shared items coded under codebook v1.2. 'Description states when to use' is reported over the 20 items the codebook's worked examples do not cover; over all 25 it reaches $\kappa = 0.779$. 'Secret documented' is not a reliability figure: the pack shows the manifest rule's verdict for every item, so both coders applied a lookup rather than exercising judgement.

Part	Field	Observed	κ
C	description states purpose	1.00	1.000
C	description names side effects	1.00	1.000
B	write capability	0.98	0.969
B	authentication mechanism	0.98	0.965
B	authentication evidence	0.98	0.964
B	requires a secret	0.98	0.964
C	confusable with a sibling	0.98	0.960
B	destructive capability	0.90	0.841
B	scope breadth	0.84	0.761
C	description explains inputs	0.82	0.725
C	description states when to use	0.95	0.643

switching Python to an AST-based extractor. Re-scoring against the same human ground truth showed recall falling to 95.3%, because the two readers have complementary blind spots rather than one dominating the other. The validated configuration is the one reported here.

4.4.1. Reliability of the human coding

Two coders independently coded the same 50 items under one shared codebook and one shared evidence pack. Agreement is reported per field rather than pooled: a pooled figure would hide which claims are safe to make.

Three things are worth reading off the table. The two fields that ask what a description *contains* — whether it states a purpose, whether it named an effect — reached complete agreement, because v1.2 reduced them to questions about the text itself. The two that require a judgement about whether a condition is *sufficient* — selection guidance and input explanation — are the lowest, at 0.643 and 0.725. And the manifest-only rule for secret declaration turned a field that had failed at 0.229 into a mechanical lookup, which is what that entry records.

The claim about selection guidance therefore stands as a point estimate. On the full 200-tool sample the coder marked 23.0% (95% CI [17.5, 29.0]); in this round the two coders

marked 8% and 12% of their 25 shared tools and agreed on 19 of the 20 items outside the worked examples. The field is reliable enough to report as a proportion, and the paper does so.

One caveat belongs in the record rather than in the table. Round 2 measured the same fields under a defective instrument: the second coder's pack carried no evidence, and two value sets were ambiguous. Round 3 is the first round in which both coders worked from the same complete material, so the two rounds are not a before-and-after of coder skill. They are a before-and-after of instrument quality, which is why the codebook revisions are described in §3.5 rather than hidden here.

5. Discussion

5.1. Where the MCP ecosystem sits against established ones

Every number in Section 4 describes a twelve-month-old ecosystem, which makes "is this a lot?" impossible to answer from the numbers alone. Placing them beside package ecosystems measured at comparable maturity changes what several of them mean.

Concentration is not unusual; the concentration mechanism is. Decan et al. compared dependency-network evolution across seven packaging ecosystems and found that growth statistics conceal very different underlying dynamics (Decan et al., 2019). Our finding that 100 of 20,540 namespaces produce 30.1% of versions is consistent with the heavy tails they report, so the concentration itself confirms the established picture rather than departing from it. What differs is where the tail comes from. In npm and PyPI, high-volume publishers are mature libraries with long dependency histories; here the top namespace is a single publisher that produced 2,387 versions of an agent's tool set. The tail is generated by automated republication of one interface, not by accumulated engineering. The practical consequence is that version counts, which are a workable activity proxy in the ecosystems Decan et al. studied, are weaker here which is what §5.2 goes on to test.

Single-release abundance contradicts an assumption carried over from package ecosystems. That 62.7% of servers publish exactly one version and never update again sits badly beside the maintenance expectations the packaging literature establishes. Abdalkareem et al. showed that npm and PyPI carry a large volume of trivial packages, so registry size overstates capability (Abdalkareem et al., 2020); our result extends that observation from package content to package trajectory. A trivial package at least gets published once and is usually a dependency of something. A single-release MCP server may never be installed by anyone, and the registry counts it exactly like a maintained one. Replication studies of technical leverage find that relying on small or inattentive maintainers carries a measurable security cost (Samaana et al., 2025); our 67.4% single-maintainer share says that most of this ecosystem is in precisely that position.

Withdrawal is rarer than in a mature registry. Li et al. found that 9.6% of Cargo packages have at least one yanked release and that the proportion kept rising (Li et al., 2023). In our snapshot, 2.1% of npm packages are deprecated and 1.1% of PyPI packages carry a yanked release. Two readings are available and we cannot separate them with cross-sectional data. Either withdrawal is genuinely rarer at this maturity, or the ecosystem is young enough that the releases which will need withdrawal have not been published yet. Li et al.'s rising trend supports the second, and the distinction matters: it implies the withdrawal rate is a lagging indicator and will rise, so a client cannot currently use low withdrawal rates as evidence of stability.

Documentation converges on the mechanical and stops there. Meng et al. observed that developers' use of API documentation varies in ways a design strategy must accommodate, and proposed guidelines accordingly (Meng et al., 2019); Fan et al. documented that documentation routinely omits the information needed to decide how to use an operation (Fan et al., 2020). Our results locate the omission precisely. Naming has converged almost completely (86.5% snake_case, enforced reverse-DNS names), and purpose statements are near-universal at 93.5%. What is missing is the judgement-bearing part: a condition for choosing the tool. This is a *sharper* version of the finding those works report, because it identifies which documentation element disappears when an interface is consumed by a model rather than read by a person. Hora et al. show that interfaces evolve faster than their documentation (Hora et al., 2018); our 93.5% ambiguity rate suggests that for agent-facing interfaces, evolution is not even required to produce the mismatch.

Credential exposure has no established baseline, so we give a mechanical one. No prior work measures credential declaration in a tool registry. The nearest frame comes from the agent-security literature. Khan et al. evaluate defences against indirect prompt injection in tool-using agents and find that restricting the available tools lowers attack success but also lowers benign utility (Khan et al., 2026). Whether that trade-off is acceptable depends on knowing what the tools can reach, which is exactly the information our RQ2 shows is not declared. Yagoubi et al. find that sensitive data leaks through tool arguments, a pathway that output-only audits do not inspect (Yagoubi et al., 2026); a client cannot audit a pathway whose inputs the manifest does not mark as sensitive. Our 34,704 unflagged credential-like variables are therefore not a cosmetic documentation gap. They are the reason two published defence strategies cannot be applied without reading source.

5.2. Enforcing the shape of a declaration is direct; enforcing the practice behind it is not

Version counts are the obvious way to describe an ecosystem's health, and this study used them throughout. We tested whether they mean what they appear to mean, against an independent signal: the releases each repository tags on

GitHub, for the 914 servers in the repository sample whose metadata resolved.

These two signals do not measure the same thing. **59.3% of those servers never tag a release at all** while continuing to publish to the registry, so their median tagged-release count is zero against a median of two registry versions. Repository releases cannot serve as ground truth for the population. For the minority that do tag releases the relationship inverts: release counts exceed registry versions for 24.9% of servers, and the registry understates releases by two times or more for 18.3%. Rank agreement is weak to moderate throughout (Spearman 0.379).

We draw a boundary rather than a conclusion, and the boundary generalises. The registry is usable for *presence* questions, such as whether a server ships updates at all and in which month it first appeared, but not for *cadence* questions. The sub-day median interval between registry versions is an ingestion artefact of batch publishing, not a release rhythm, and we do not present it as one.

This is one instance of a more general property, and it is the claim we would like other registry designers to take from this paper:

A registry can enforce the *shape* of a declaration directly, because shape is checkable at the moment of submission. Enforcing the *practice* behind the declaration takes post-hoc mechanisms, since practice happens afterwards. Those mechanisms exist, and they are more expensive and slower, so a registry that enforces only shape should be read as making a choice rather than facing a limit. The distance between the two is measurable, and measuring it is the first thing a population study of a registry should do.

The evidence here is unusually clean because the same entity enforces both things at different strengths. The MCP registry enforces naming: no server in 35,388 omits the slash separator, none departs from a reverse-DNS namespace, and none puts a space in the server segment. It enforces nothing about release practice, and on that axis the registry and the repositories disagree for 89.1% of servers in the sample. The registry also declares credential semantics without enforcing them: a publisher may mark a variable secret, and 19.9% of credential-like variables are left unmarked. Shape is a submission-time property; practice and semantics are not, and the gap between them is where every measurement error in this paper originates.

Two consequences follow for anyone building a registry of agent-facing artefacts. First, aggregate activity metrics computed from a registry inherit its blind spots, so a census should test at least one headline series against an independent source before reporting trends, which is what §5.2 does here. Second, properties worth enforcing are exactly those that are shape-like: presence of a repository link, presence of a licence, presence of a flag. Properties requiring judgement, such as whether a capability is broad or whether a description

is sufficient, resist enforcement, and the next subsection explains why.

5.3. What the agreement pattern shows

Table 8 was intended as a reliability disclosure and turned into the paper's most transferable result. The fields do not fail uniformly. They divide cleanly by what the coder is asked to judge.

Fields that ask whether a feature is *present* reached moderate agreement or better: whether a description states a purpose (1.000), whether it named a side effect (1.000), whether a server can write (0.969), which authentication mechanism it uses (0.965), whether it requires a secret (0.964), whether a server can destroy (0.841), how broad its scope is (0.761). Coders looking for the same visible thing find it, and two readers converge.

Fields that ask whether something is *absent*, *sufficient* or *implied* collapse to the lowest values in the table: whether a description states when to choose the tool (0.643) and whether it explains its inputs (0.725). These are not harder to read. They are harder to agree on, because each requires a coder to decide what the absence of a statement means. Under round 2's codebook the two coders on selection guidance marked 6 and 1 of the same 25 items; tightening the rule and adding adjudicated examples moved them to 0.643, and the residual disagreements stay of the same kind: whether a stated purpose implies a selection condition, or whether an explicit condition, trigger or ordering rule is required.

That mechanism explains the ecosystem's own behaviour, and this is the part we consider new. A registry can check a presence at submission; a publisher can supply one cheaply. Both sides therefore converge on the properties that are presence-like, which is why naming is near-uniform and purpose statements are near-universal. Sufficiency has no such equilibrium. There is no submission-time check for it, and no shared standard for it: our own coders, working from the same rule and the same evidence, still agree least on exactly those fields. The ecosystem is not under-documented because publishers are careless. It is documented exactly to the depth that a registry can reward, and no further.

Three practical consequences follow. A study that reports a single pooled κ will average a reliable measurement with an unreliable one and learn nothing about either; per-field reporting is not pedantry. A governance proposal that requires publishers to document sufficiency will not converge, because the requirement cannot be adjudicated. And the parts of an agent interface that most affect whether a model selects the right tool are precisely the parts that resist both enforcement and measurement. That is a statement about the limits of documentation-based governance rather than about this registry.

5.4. How the two propositions could be tested elsewhere

Both propositions were derived from one registry during its first year, so the useful question is not whether they are

attractive but what would count as evidence against them. Each is stated below with the observation that would falsify it and the data a follow-up would need. Neither test requires access we did not have.

P1: enforcement strength at submission does not predict agreement with independent practice. What would falsify it is a registry that enforces only shape and whose version counts nevertheless track an independent release signal closely. Running that test requires only what is already here: for any registry with a published version history and a linked source repository, compute rank agreement between the two, as in §5.2. A Spearman correlation above roughly 0.8 would show that the gap we measured is a property of this registry or of its age rather than of submission-time enforcement.

A stronger design compares registries. A useful version takes two catalogues that differ in enforcement and are otherwise comparable, for example one that verifies publisher identity and manifest fields at submission and one that does not, and measures the same quantity in both. Chrome Web Store listings and VS Code Marketplace extensions are candidates: both publish version histories, both link repositories, and both differ from the MCP registry in that they run post-hoc review as well as submission-time checks. That difference is the point of the design, because it separates the two clauses of P1. If post-hoc review narrows the gap while submission-time checks do not, the proposition gains the causal reading we have avoided giving it here.

P2: coding agreement tracks whether a field asks about presence or about sufficiency. Here falsification would be a field set in which sufficiency judgements reach the agreement that presence judgements reach. Again the test is a replication, this time with the codebook itself: apply §3.5's fields to another catalogue of declared interfaces, code a sample with two independent coders, and compare the per-field κ ordering. Model cards on Hugging Face, extension manifests in the VS Code Marketplace and plugin descriptions in OpenAI's tool ecosystem are all candidates, and all publish the same kinds of artefact: a name, a free-text description, a machine-readable parameter specification, and an optional link to source.

That prediction is specific enough to fail. We expect purpose and input declaration to reach moderate agreement or better and selection guidance to stay below 0.40, in any of those catalogues, because the asymmetry we attribute to enforcement is really about what a coder can decide without inference. If selection guidance instead reaches substantial agreement elsewhere, the explanation in §5.3 is wrong and the mechanism lies in the MCP artefacts rather than in the coding task.

Three limitations of these designs should be stated, because they bound what a positive result would mean. Registry enforcement strength is not a manipulable variable, so a cross-registry comparison is observational and confounded by age, audience and governance model. Coding agreement depends on the coder population as well as the field, so a replication with different coders confounds the two unless

the same codebook and evidence-pack procedure are used, which is why both ship with this paper. And neither test addresses the case that matters most for practice: whether a client that reads only the declared artefacts makes better tool choices than one that does not. That is an experiment on agents, not on registries, and it is the natural next step.

5.5. What would have to become machine-checkable

Our results suggest a short list of properties that are consequential, cheap to declare and currently absent or optional. Each already has a natural home in the manifest; none requires a new protocol.

1. **A required secret flag on every credential-like variable.** A client cannot redact what the manifest does not mark, and 19.9% of declared variables are in that position today.
2. **A declared capability class.** Whether a server writes, deletes or spends, and tool definitions answer it, but no artefact a client can read before installing states it. 22.0% of coded servers combine broad scope with write capability, and a client currently learns this only by reading code.
3. **A selection condition on tool descriptions.** Descriptions state purpose 93.5% of the time but a condition for choosing the tool only 23.0% of the time, in an ecosystem where 93.5% of tools have a confusable sibling.
4. **A reachability declaration for auditing.** 23.3% of servers link no repository. Making the absence explicit would let a client distinguish "closed source" from "unfinished manifest".

5.6. Implications for practice

An engineer deciding today whether to install an MCP server can check three things and no more: whether a package exists and is licensed (30.1% of npm and 85.0% of PyPI packages are not), how many versions have been published (with the caveat above), and which environment variables a manifest declares and whether it flags them secret. Everything else requires reading source: capability, scope, and whether the description distinguishes one tool from another. For a quarter of servers there is no source to read.

For a publisher the same list inverts into inexpensive improvements, and the ambiguity result suggests the change of highest value is disambiguating siblings rather than writing better prose.

6. Threats to Validity

We follow the four-category scheme and mark which threats we tested rather than merely stated.

6.1. Construct validity

Registry timestamps are not release dates. Tested, and reported in §5.2. The registry records ingestion; 59.3% of

servers never tag a release at all. We restrict timing claims to presence and rephrase every cadence claim accordingly.

Extraction coverage as a proxy for documentation quality. A server whose tools cannot be extracted may generate them at runtime or document them elsewhere. Recall was established against human ground truth on 150 packages and is reported beside every tool-level claim.

Secret flagging is a publisher declaration. A variable not marked secret is not necessarily mishandled at runtime. Our claim is narrower and exact: a client cannot *automatically* redact what the manifest does not mark.

The input-schema figure is a lower bound. Only 3.2% of tools expose a schema our parser recovers. Publishers often pass schemas as variables or generate them at runtime, so this measures static readability rather than absence.

6.2. Internal validity

Registry-side validation filters. Absence of plaintext, bare-IP and localhost endpoints is partly a property of the registry's validation, not of publisher discipline. We report it as evidence about the gatekeeper.

Version-count concentration. A handful of namespaces publish thousands of versions, so outliers dominate means and totals. We report every mean beside a median and quote the concentration table wherever activity is discussed.

Codebook defects discovered during coding. The first inter-rater round produced near-zero agreement on every field. Investigation showed two causes: the handoff to the second coder carried no evidence, and two fields had defective value sets. We revised the codebook twice, rebuilt the evidence pack, and ran a further round in which both coders worked from the same complete material. Only that round is reported as reliability.

Two limits on it belong here. The pack supplies the manifest rule's verdict for the secret-declaration field, so that entry measures rule application rather than judgement; and five of the coded tools appear as worked examples in the codebook, so the selection-guidance figure is computed over the 20 items those examples do not cover. Both are stated in Table 8 rather than only here.

Sampling drift. Any study of this population contends with growth. We freeze the harvest at one date, record its digest, and draw every sample from that snapshot.

6.3. External validity

One registry, one protocol version. Coverage is one registry during its first year, with 92.3% of versions on a single manifest schema. Findings may not transfer to other agent-tool conventions.

Two package registries. Tool extraction covers only npm and PyPI, and only the 1,200-package sample. The 10.9% of package declarations in the MCP bundle format and the 9.1% in OCI images are not covered.

Rapid change. The ecosystem is twelve months old at the freeze. Conclusions are properties of the window analysed; the pipeline is built to be re-run at a later freeze.

6.4. Conclusion validity

Recall and precision are estimated from counts, not identities. The coder recorded how many tools each package exposes rather than naming them, so we estimate matching as the per-package minimum of the two counts. That estimate holds in aggregate but cannot detect a package that over- and under-counts at the same time.

Cross-source comparisons are descriptive. We report no *p*-values; the relationships in §4 are associations.

6.5. Reproducibility

Sources are mutable and packages can be unpublished. Raw responses are cached, every generated table is written to disk, and a manifest of SHA-256 digests covers the snapshot, the samples and the outputs.

7. Conclusion

We measured the MCP server ecosystem at a single frozen point, from the registry outward, and validated each automated step against human judgement. The ecosystem is large and growing steeply: 114,604 published versions across 35,388 servers, with 27,863 versions in the final month alone. It is also unevenly built. Most servers publish once and stop; 100 namespaces produce 30.1% of all versions; 23.3% of servers cannot be traced to source at all; and 19.9% of declared environment variables look like credentials without being flagged as such.

The result we did not expect concerns how the ecosystem documents itself. Its conventions have converged on the mechanical parts, namely one naming style, one manifest schema and near-universal description coverage, and have left the judgement-dependent parts behind. Descriptions say what tools do; they rarely say when to choose one, and 93.5% of tools have a sibling an agent could confuse them with. Two human coders could only reach fair agreement about which descriptions give selection guidance, which suggests the difficulty is not writerly but conceptual, and it will not yield to asking publishers to write more.

The practical consequence is a short agenda: make the secret flag mandatory, declare a capability class, add a selection condition, and state honestly when no source is available. None requires a protocol change. All four would move the agent-tool layer from partly auditable to mostly auditable.

Two findings travel beyond this protocol. A census of a registry measures publication rather than practice, and here the two disagree for 89.1% of the servers we could check, so any population study of a young registry should test at least one headline series against an outside source before reporting a trend. And an ecosystem documents itself to the depth its registry rewards: fields that ask whether a property is present are cheap to supply and easy to agree on, so they converge, while fields that ask whether something is sufficient are neither, so they do not. Both statements are consequences of measuring the artefacts rather than the systems they describe, and both extend to any registry of agent-facing components.

CRedit authorship contribution statement

Junan Chen: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Visualization.

Declaration of competing interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Data availability

The replication package is deposited in Science Data Bank under DOI 10.57760/sciencedb.013wx and cited in the reference list (Chen, 2026). It contains the collection and analysis pipeline, the seeded samples, the generated result tables, the coding codebooks with their worksheets, and a manifest of SHA-256 digests covering every artefact. Running the pipeline regenerates all raw data from the public sources listed in Table 1. The raw registry snapshot is not redistributed, because the upstream registry is its authoritative copy; the manifest lets the snapshot used here be verified against a fresh harvest.

Declaration of generative AI and AI-assisted technologies

Generative AI tools were used to help write and refactor the collection and analysis code, and to assist with language editing of the manuscript. The author reviewed and verified all code, all generated tables and all text, and takes full responsibility for the content.

Acknowledgements

The author thanks the second coder, who contributed the independent coding reported in Section 4.4.1 under a codebook shared with the author, and two colleagues who read an earlier draft.

References

- Abdalkareem, R., Oda, V., Mujahid, S., Shihab, E., 2020. On the impact of using trivial packages: an empirical case study on npm and pypi. *Empirical Software Engineering* 25, 1168–1204. doi:10.1007/s10664-019-09792-9.
- Ball, B.L., Long, N., Fleming, K., Goldwasser, D., 2026. OpenStudio-MCP: a model context protocol (MCP) server for AI agent-driven building energy modelling. *SoftwareX* doi:10.1016/j.softx.2026.103020.
- Chen, J., 2026. Replication package for: The Model Context Protocol server ecosystem: an empirical study of quality, security and evolution. URL: <https://doi.org/10.57760/sciencedb.013wx>, doi:10.57760/sciencedb.013wx.
- Cogo, F.R., Oliva, G.A., Bezemer, C.P., Hassan, A.E., 2021a. An empirical study of same-day releases of popular packages in the npm ecosystem. *Empirical Software Engineering* 26. doi:10.1007/s10664-021-09980-6.
- Cogo, F.R., Oliva, G.A., Hassan, A.E., 2021b. An empirical study of dependency downgrades in the npm ecosystem. *IEEE Transactions on Software Engineering* 47, 2457–2470. doi:10.1109/TSE.2019.2952130.
- Cohen, J., 1960. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement* 20, 37–46. doi:10.1177/001316446002000104.
- Decan, A., Mens, T., Grosjean, P., 2019. An empirical comparison of dependency network evolution in seven software packaging ecosystems. *Empirical Software Engineering* 24, 381–416. doi:10.1007/s10664-017-9589-y.
- Fan, Q., Yu, Y., Wang, T., Yin, G., Wang, H., 2020. Why API documentation is insufficient for developers: an empirical study. *Science China Information Sciences* 63. doi:10.1007/s11432-019-9880-8.
- Hora, A., Robbes, R., Valente, M.T., Anquetil, N., Etien, A., Ducasse, S., 2018. How do developers react to API evolution? A large-scale empirical study. *Software Quality Journal* 26, 191–227. doi:10.1007/s11219-016-9344-4.
- Khan, A., AlKhanbashi, K., Mohamed, A., 2026. Evaluating indirect prompt injection defenses in tool-using LLM agents: security, utility and cost. *Computers* 15. doi:10.3390/computers15090570.
- Landis, J.R., Koch, G.G., 1977. The measurement of observer agreement for categorical data. *Biometrics* 33, 159–174. doi:10.2307/2529310.
- Li, H., Cogo, F.R., Bezemer, C.P., 2023. An empirical study of yanked releases in the Rust package registry. *IEEE Transactions on Software Engineering* 49. doi:10.1109/TSE.2022.3152148.
- Li, X., Moreschini, S., Pecorelli, F., Taibi, D., 2022. OSSARA: abandonment risk assessment for embedded open source components. *IEEE Software* 39. doi:10.1109/MS.2022.3163011.
- Liu, Y., Lo, S.K., Lu, Q., Zhu, L., Zhao, D., Xu, X., Harrer, S., Whittle, J., 2025a. Agent design pattern catalogue: A collection of architectural patterns for foundation model based agents. *Journal of Systems and Software* 220, 112278. doi:10.1016/j.jss.2024.112278.
- Liu, Y., Tiwari, D., Bogdan, C., Baudry, B., 2025b. Detecting and removing bloated dependencies in commonjs packages. *Journal of Systems and Software* 230, 112509. doi:10.1016/j.jss.2025.112509.
- Meng, M., Steinhart, S., Schubert, A., 2019. How developers use API documentation. *Communication Design Quarterly* 7. doi:10.1145/3358931.3358937.
- Model Context Protocol, a. Model Context Protocol specification. <https://modelcontextprotocol.io>. Accessed 24 September 2026.
- Model Context Protocol, b. Official MCP registry API. <https://registry.modelcontextprotocol.io>. Accessed 24 September 2026; snapshot frozen at this date.
- Myers, B.A., Stylos, J., 2016. Improving API usability. *Communications of the ACM* 59, 62–69. doi:10.1145/2896587.
- Nocera, S., Romano, S., Di Penta, M., Francese, R., Scanniello, G., 2025. On the adoption of software bill of materials in open-source software projects. *Journal of Systems and Software* 230, 112540. doi:10.1016/j.jss.2025.112540.
- Othmani, T., Ben Yahia, S., Lalaguna, A., 2026. Model context protocol-based agentic react large language model for adaptive traffic signal control. *Future Generation Computer Systems* doi:10.1016/j.future.2025.108339.
- Qi, Q., Cao, J., 2024. An empirical study on downstream dependency package groups in software packaging ecosystems. *IET Software* 2024. doi:10.1049/2024/4488412.
- Robillard, M.P., 2009. What makes APIs hard to learn? Answers from developers. *IEEE Software* 26, 27–34. doi:10.1109/MS.2009.193.
- Rosales, E., Basso, M., Rosà, A., Binder, W., 2023. Large-scale characterization of Java streams. *Software: Practice and Experience* 53, 1763–1792. doi:10.1002/spe.3213.
- Samaana, H., Costa, D.E., Abdellatif, A., Shihab, E., 2025. Opportunities and security risks of technical leverage: a replication study on the npm ecosystem. *Empirical Software Engineering* 30. doi:10.1007/s10664-025-10648-8.

- Umar Zeshan, M., Ibiyo, M., Di Sipio, C., Nguyen, P.T., Di Ruscio, D., 2026. Many hands make light work: An LLM-based multi-agent system for detecting malicious PyPI packages. *Journal of Systems and Software* 236, 112792. doi:10.1016/j.jss.2026.112792.
- Venigalla, A.S.M., Chimalakonda, S., 2025. Is there a correlation between readme content and project meta-characteristics? *Software: Practice and Experience* 55, 589–609. doi:10.1002/spe.3390.
- Weeraddana, N.R., Alfadel, M., McIntosh, S., 2024. Dependency-induced waste in continuous integration: an empirical study of unused dependencies, in: *Proceedings of the ACM on Software Engineering*. doi:10.1145/3660823.
- Yagoubi, F.E., Badu-Marfo, G., Mallah, R.A., 2026. AgentLeak: a benchmark for internal-channel privacy leakage in multi-agent LLM systems. *IEEE Access* doi:10.1109/ACCESS.2026.3704541.
- Zhou, J., Chen, J., Wang, M., Dong, B., Zhou, X., 2026. Large language model orchestrated workflow integrating retrieval-augmented generation and model context protocol. *Building and Environment* doi:10.1016/j.buildenv.2026.114553.

Junan Chen is a researcher at the University of Electronic Science and Technology of China. His work focuses on empirical software engineering for agent-facing infrastructure, with an emphasis on measuring software ecosystems from published artefacts and on making such measurement reproducible.




Figure 1: Credential-like environment variables, split by whether the manifest flags them secret. A client cannot redact automatically what the manifest does not mark. Source: `t11_credential_hygiene.csv`.

The MCP server ecosystem




Figure 2: Ecosystem growth over the twelve months of the snapshot. Bars give new versions published per month; the line gives cumulative distinct servers. Both are registry counts, so they measure publication rather than deployment. Source: `t07_publishing_timeline.csv`.




Figure 3: Published versions per server. The distribution is dominated by single-release servers, which is why every activity statistic in this paper carries a median beside its mean. Source: `t08_versions_per_server.csv`.